

Julia ODEInterface Experiments

Vishal Sontakke

Contents

1	Non-stiff equations	2
1.1	The restricted three body problem (AREN)	2
1.2	Movement of hanging rope (ROPE)	3
1.3	Code Performance	4
1.3.1	AREN	4
1.3.2	ROPE	5
2	Stiff equations	6
2.1	Van der Pol oscillator (VDPOL)	6
2.2	Robertson Chemical Kinetics (ROBER)	7
2.3	Code performance	8
2.3.1	VDPOL	8
2.3.2	ROBER	10
3	Appendix of Codes	12
3.1	savePlotPNG	12
3.2	ArenPrecisionTest	13
3.3	RopePrecisionTest	15
3.4	vdpolPrecisionTest	17
3.5	RoberPrecisionTest	20

1 Non-stiff equations

The problems chosen for our tests are the following:

1.1 The restricted three body problem (AREN)

We will first look at the restricted three body problem. One considers two bodies of masses $1 - \mu$ and μ in circular rotation in a plane and a third body of negligible mass moving around in the same plane. The equations are (see e.g., the classical textbook Szebehely 1967):

$$\begin{aligned} y_1'' &= y_1 + 2y_2' - \mu' \frac{y_1 + \mu}{R_1} - \mu \frac{y_1 - \mu'}{R_2}, \\ y_2'' &= y_2 - 2y_1' - \mu' \frac{y_2}{R_1} - \mu \frac{y_2}{R_2}, \\ R_1 &= ((y_1 + \mu)^2 + y_2^2)^{3/2}, \quad R_2 = ((y_1 - \mu')^2 + y_2^2)^{3/2}, \\ \mu &= 0.012277471, \quad \mu' = 1 - \mu. \end{aligned} \tag{1.1}$$

There exist initial values, for example:

$$\begin{aligned} y_1(0) &= 0.994, \quad y_1'(0) = 0, \quad y_2(0) = 0, \\ y_2'(0) &= 2.00158510637908252240537862224, \\ t_{end} &= 17.0652165601579625588917206249, \end{aligned} \tag{1.2}$$

such that the solution is periodic with period t_{end} .




Figure 1.1: An Arenstorf Orbit computed using Dormand & Prince

1.2 Movement of hanging rope (ROPE)

We look at the movement of a hanging rope of length 1 under gravitation and under the influence of a horizontal force

$$F_y(t) = \left(\frac{1}{\cosh(4t - 2.5)} \right)^4 \quad (1.3)$$

acting at the point $s = 0.75$ as well as a vertical force

$$F_x(t) = 0.4 \quad (1.4)$$

acting at the endpoint $s = 1$.

If this problem is discretized, then Lagrange theory leads to the following equations for the unknown angles θ_k :

$$\begin{aligned} \sum_{k=1}^n a_{lk} \ddot{\theta}_k = & - \sum_{k=1}^n b_{lk} \dot{\theta}_k^2 - n \left(n + \frac{1}{2} - l \right) \sin \theta_l \\ & - n^2 \sin \theta_l \cdot F_x(t) + \begin{cases} n^2 \cos \theta_l \cdot F_y(t) & \text{if } l \leq 3n/4 \\ 0 & \text{if } l > 3n/4, \end{cases} \quad l = 1, \dots, n \end{aligned} \quad (1.5)$$

where

$$a_{lk} = g_{lk} \cos(\theta_l - \theta_k), \quad b_{lk} = g_{lk} \sin(\theta_l - \theta_k), \quad g_{lk} = n + \frac{1}{2} - \max(l, k). \quad (1.6)$$

We choose

$$n = 40, \quad \theta_l(0) = \dot{\theta}_l(0) = 0, \quad 0 \leq t \leq 3.723. \quad (1.7)$$




Figure 1.2: Movement of the rope with respect to time

1.3 Code Performance

The three available non-stiff solvers were applied to the above mentioned problems with $Tol = 10^{-3}$, $Tol = 10^{-3-1/8}$, $Tol = 10^{-3-2/8}$, $Tol = 10^{-3-3/8}$, ... up to $Tol = 10^{-14}$, then the numerical result at the output points were compared with a "reference solution" using the infinity norm. The "reference solution" is computed using $Tol = 1.0 \times 10^{-16}$, since quadruple precision is not available on the current hardware (as compared to [?]). The underlying Fortran codes would have to be tweaked to get quadruple precision using software since Julia can emulate Arbitrary Precision floating point numbers using GNU Multiple Precision Arithmetic Library (GMP) and the GNU MPFR Library.

1.3.1 AREN




Figure 1.3: Precision vs. Number of function evaluations (as computed on Julia using ODEInterface)

1.3.2 ROPE

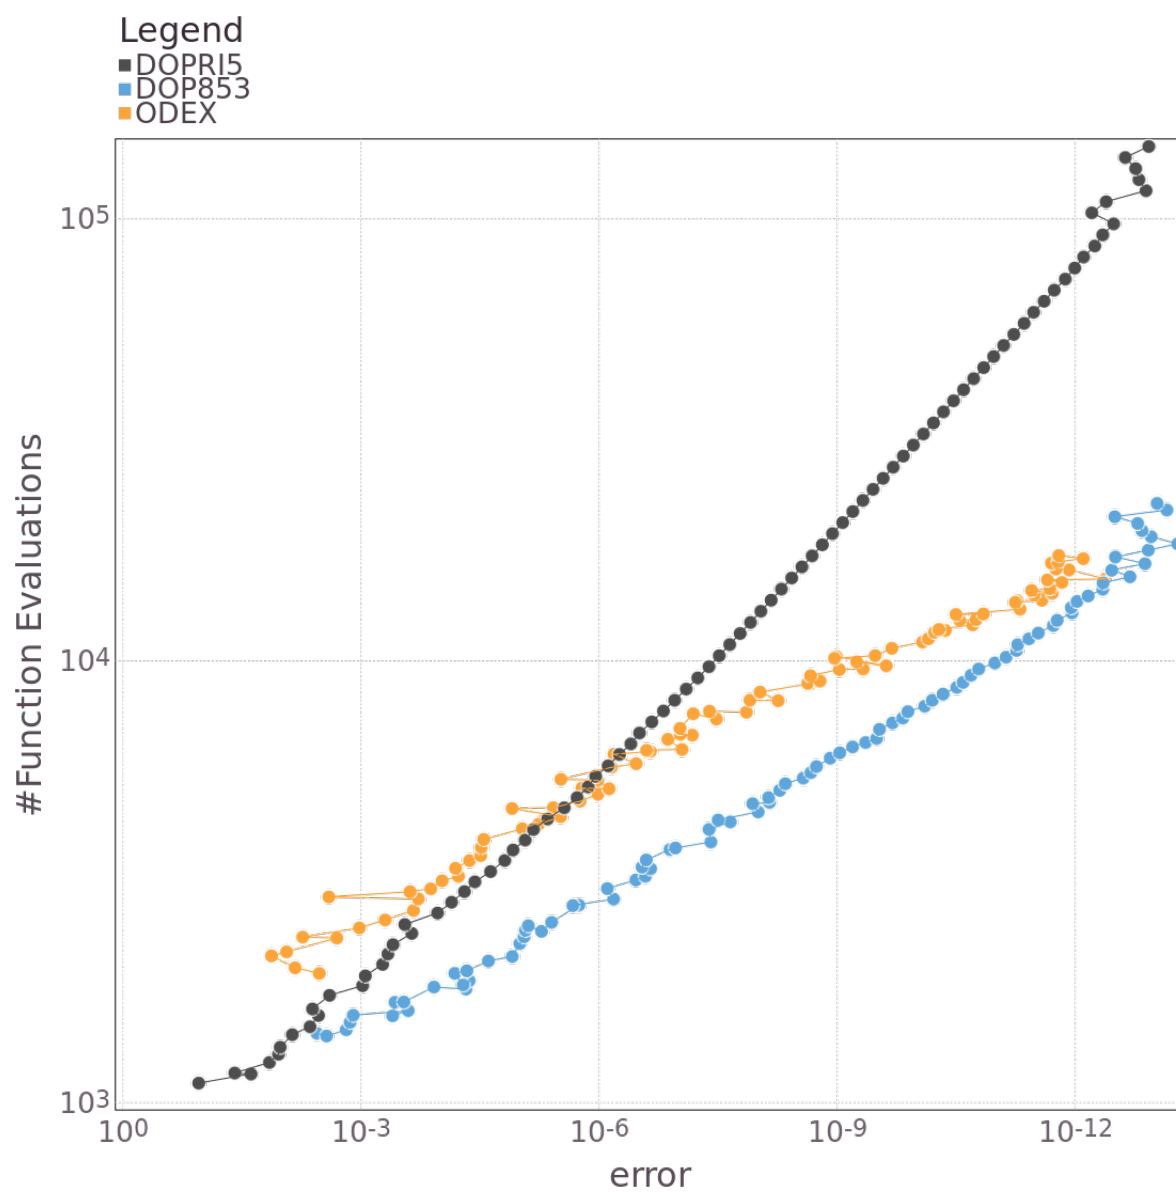


Figure 1.4: Precision vs. Number of function evaluations (as computed on Julia using ODEInterface)

2 Stiff equations

The problems chosen for our tests are the following:

2.1 Van der Pol oscillator (VDPOL)

In dynamics, the Van der Pol oscillator is a non-conservative oscillator with non-linear damping. It evolves in time according to the second-order differential equation:

$$\frac{d^2x}{dt^2} - \mu (1 - x^2) \frac{dx}{dt} + x = 0, \quad (2.1)$$

where x is the position coordinate-which is a function of time t , and μ is a scalar parameter indicating the non-linearity and the strength of the damping.

The above equation can be converted into a system of first-order differential equations as follows:

$$\begin{aligned} x_1' &= x_2 \\ x_2' &= ((1 - x_1^2) x_2 - x_1) / \epsilon, \quad \epsilon = 10^{-6} \end{aligned} \quad (2.2)$$

We perform a rescaling as follows:

$$\begin{aligned} \tilde{t} &= t/\mu, \quad x_1(\tilde{t}) = x(t), \quad x_2(\tilde{t}) = \mu \frac{dx}{dt}(t) \\ \text{and we set} \quad \frac{1}{\mu^2} &= \epsilon \end{aligned} \quad (2.3)$$

This transformation makes the steady state approximation independent of μ .

The initial conditions for this problem are:

$$\begin{aligned} x_1(0) &= 2, & x_2(0) &= 0 \\ t_{out} &= 1, 2, 3, 4, \dots, 11. \end{aligned} \quad (2.4)$$

The times t_{out} will be the points at which the solution will be taken for comparison.

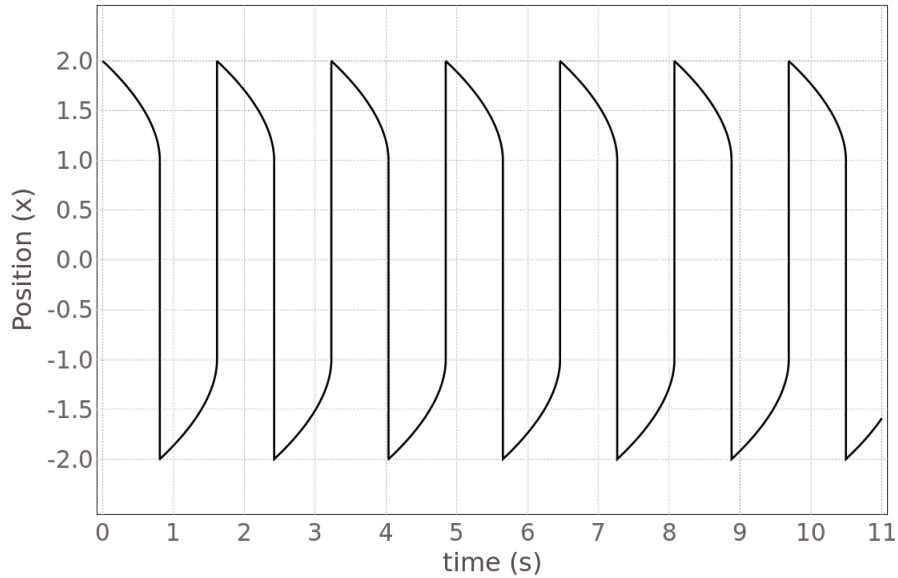


Figure 2.1: Solution of equation (??) computed using SEULEX solver

2.2 Robertson Chemical Kinetics (ROBER)

The problem describes the kinetics of an auto-catalytic reaction given by Robertson as follows:

$$\begin{aligned}x_1' &= -0.04x_1 + 10^4x_2x_3 \\x_2' &= 0.04x_1 - 10^4x_2x_3 - 3 \cdot 10^7x_2^2 \\x_3' &= 3 \cdot 10^7x_2^2.\end{aligned}\tag{2.5}$$

with initial conditions:

$$x_1(0) = 2, \quad x_2(0) = 0, \quad x_3(0) = 0.\tag{2.6}$$

one of the most prominent examples of the "stiff" literature. It was usually treated on the interval $0 \leq t \leq 40$, until Hindmarsh discovered that many codes fail if t becomes very large (10^{11} say). The reason is that whenever the numerical solution of x_2 accidentally becomes negative, it then tends to $-\infty$ and the run ends by overflow. We have therefore chosen $t_{out} = 1, 10, 10^2, 10^3, \dots, 10^{11}$.




Figure 2.2: Solution of equation (??) using SEULEX solver

2.3 Code performance

The three available stiff solvers were applied to the above mentioned problems with $Tol = 10^{-2}, Tol = 10^{-2-1/4}, Tol = 10^{-2-2/4}, Tol = 10^{-2-3/4}, \dots$ up to $Tol = 10^{-10}$.

We set the relative error tolerance to be $RTol = Tol$ and the absolute error tolerance $ATol = 10^{-6} \cdot Tol$ for the problem ROBER and $ATol = Tol$ for VDPOL. Then the numerical results were compared with a "reference solution" using the infinity norm (the norm is taken over all components and all output points). The "reference solution" is obtained from <http://www.unige.ch/~hairer/testset/testset.html>.

2.3.1 VDPOL

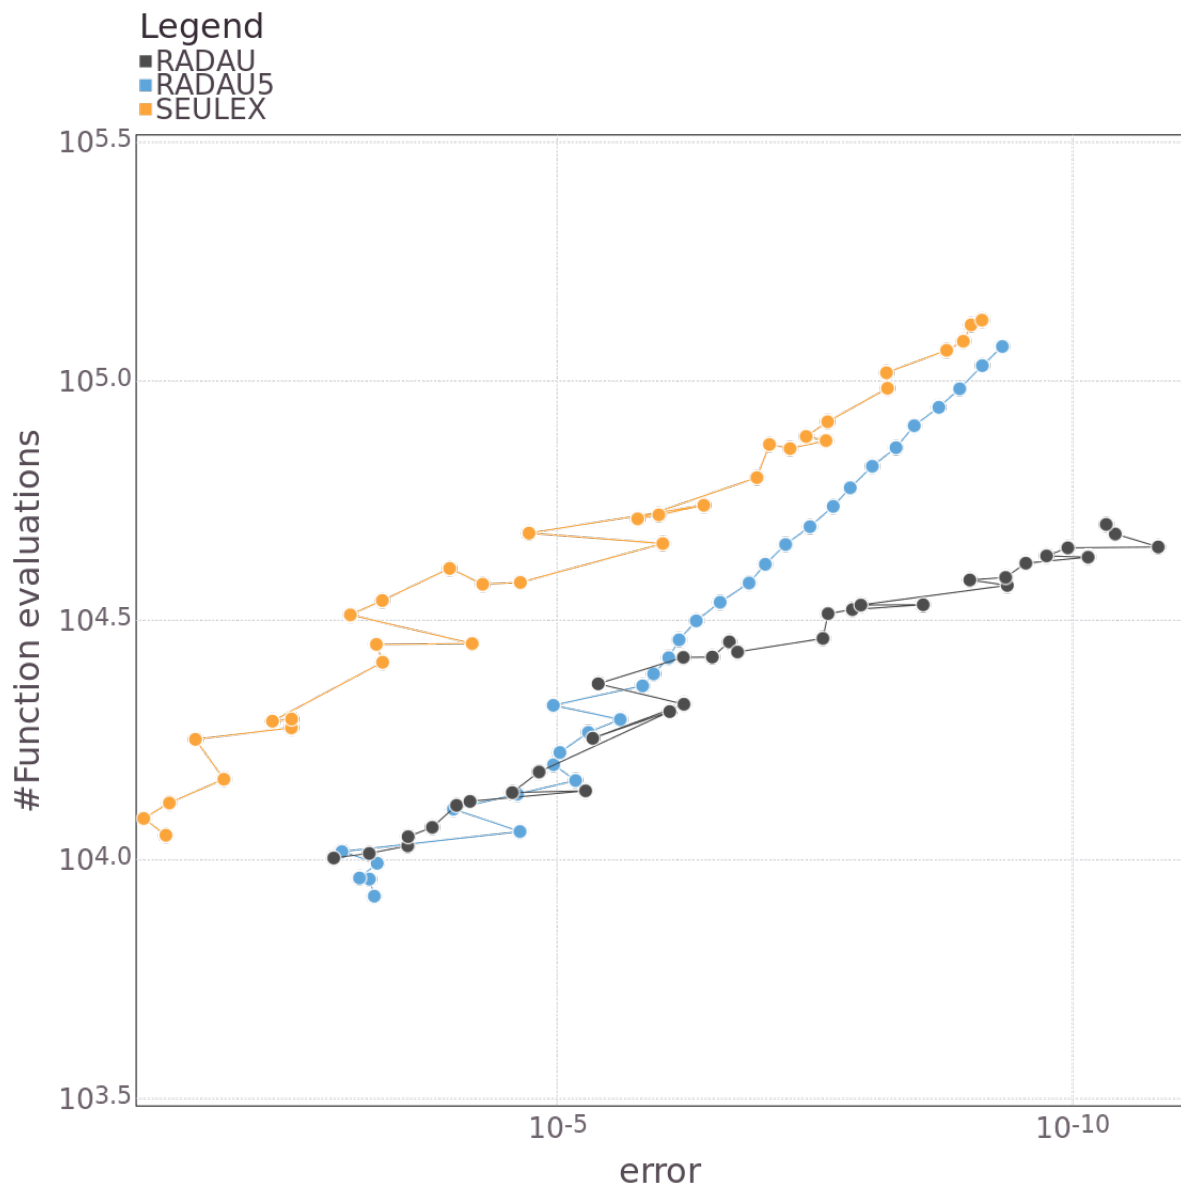


Figure 2.3: Precision vs. Number of function evaluations (as computed on Julia using ODEInterface)




Figure 2.4: Precision vs. Floating point operations (as computed on Julia using ODEInterface)

The following factors were chosen for the various operations:

- Right-hand side function calls = 5 flops
- LU Decomposition = $\lceil \frac{2}{3} \times 8 \rceil = 6$ flops
- Forward or Backward Substitution = 4 flops
- Jacobian computation using Automatic Differentiation = $\lceil 1.5 * 5 \rceil = 8$ flops

2.3.2 ROBER

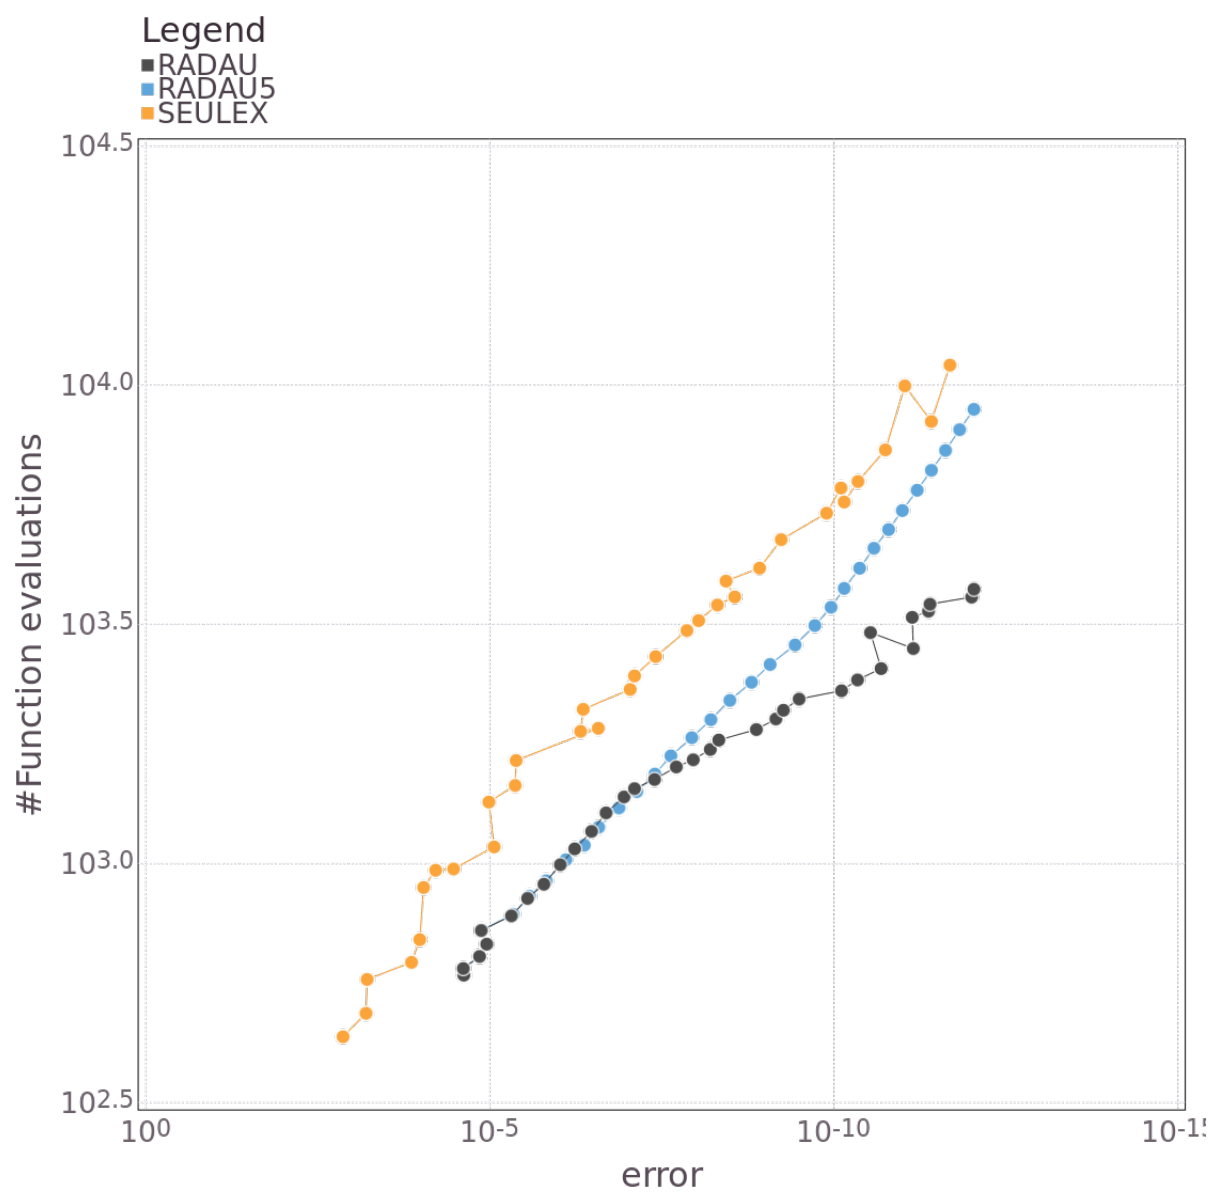


Figure 2.5: Precision vs. Number of function evaluations (as computed on Julia using ODEInterface)




Figure 2.6: Precision vs. Floating point operations (as computed on Julia using ODEInterface)

The following factors were chosen for the various operations:

- Right-hand side function calls = 5 flops
- LU Decomposition = $\lceil \frac{2}{3} \times 8 \rceil = 6$ flops
- Forward or Backward Substitution = 4 flops
- Jacobian computation using Automatic Differentiation = $\lceil 1.5 * 5 \rceil = 8$ flops

3 Appendix of Codes

The codes over here are from March,2016. For updated codes refer <https://github.com/sonVishal/ODEInterface.jl/tree/master/myExamples/Codes>

3.1 savePlotPNG

```
using Gadfly
using Colors

##### Function for saving plots #####
# Input:
# fileName = Name of the file where the plot is to be stored
#           (with or without extension)
# f_e = Array containing function evaluations as columns for each solver
# err = Array containing erros as columns for each solver
# solverNames = Array containing the names of solvers used in respective order
# plotSize = size of the plot to be created
#
# Values have been tuned for a graph similar to the one in
# Solving Ordinary Differential Equations I by
# Hairer, Ernst, Nørsett, Syvert P., Wanner, Gerhard
# page: 252
#####
function savePlotPNG(fileName,f_e,err,solverNames,
    plotSize=[30cm,30cm])

    numOfLayers = length(solverNames);

    if !contains(fileName, ".")
        fileName = string(fileName, ".png");
    end

    plotColorsHex = ["#4D4D4D", "#5DA5DA", "#FAA43A", "#60BD68",
        "#F17CB0", "#B2912F", "#B276B2", "#DEC3F3", "#F15854"];
    plotColors = [parse(Colorant,c) for c in plotColorsHex];

    majorFontSize = 24pt;
    minorFontSize = 20pt;
    pointSize = 5pt;

    myplot = plot(Scale.x_log10,Scale.y_log10,
        Coord.cartesian(xflip=true),
        Guide.manual_color_key("Legend",solverNames,plotColorsHex[1:numOfLayers]),
        Guide.xlabel("error"),Guide.ylabel("#Function Evaluations"),
        Theme(major_label_font_size=majorFontSize,panel_stroke=colorant"black",
            minor_label_font_size=minorFontSize,key_title_font_size=majorFontSize,
            key_label_font_size=minorFontSize,key_position=:top,key_max_columns=1));

    for i = 1:numOfLayers
        push!(myplot,layer(x=err[:,i],y=f_e[:,i],Geom.point,Geom.path,
            Theme(default_color=plotColors[i],default_point_size=pointSize)));
    end

    draw(PNG(fileName,plotSize[1],plotSize[2]),myplot)
    return nothing
end
```

3.2 ArenPrecisionTest

```
# Load all the required packages
using ODEInterface
@ODEInterface.import_huge
loadODESolvers();

# Define the system of ODEs
function threebody(t,x,dx)
    mu = 0.012277471; ms = 1 - mu;
    r1 = vecnorm(x[1:2]-[-mu,0]);
    r2 = vecnorm(x[1:2]-[ ms,0]);

    dx[1] = x[3];
    dx[2] = x[4];
    dx[3] = x[1] + 2*x[4] - ms*(x[1]+mu)/r1^3 - mu*(x[1]-ms)/r2^3;
    dx[4] = x[2] - 2*x[3] - ms*      x[2]/r1^3 - mu*      x[2]/r2^3;

    return nothing
end

# Flag to check if all solvers were successful
printFlag = true;

# Initial conditions
t0 = 0.0; T = 17.06521656015796; x0=[0.994, 0.0, 0.0, -2.001585106379082];

# Get "reference solution"
opt = OptionsODE(OPT_EPS => 1.11e-16,OPT_RHS_CALLMODE => RHS_CALL_INSITU,
OPT_RTOL => 1e-16,OPT_ATOL=>1e-16);
(t,x_ref,retcode,stats) = dop853(threebody,t0, T, x0, opt);

if retcode != 1
    println("Reference solution failed")
else
    # Initialization for the loop
    # f_e = function evaluations
    f_e = zeros(Int32,89,3);
    # err = error for last step using infinity norm
    err = zeros(Float64,89,3);

    # solverNames = names of the solvers used for the plot
    solverNames = ["DOPRI5","DOP853","ODEX"];

    # Compute all the solutions
    for i=0:88

        # Set up the tolerance
        Tol = 10^(-3-i/8);

        # Set up solver options
        opt = OptionsODE(OPT_EPS => 1.11e-16,OPT_RHS_CALLMODE => RHS_CALL_INSITU,
OPT_RTOL => Tol,OPT_ATOL => Tol);

        # Solve using DOPRI5
        (t,x,retcode,stats) = dopri5(threebody,t0, T, x0, opt);
```

```

    # Check if solver was successful
    if retcode != 1
        printFlag = false;
        break;
    end
    f_e[i+1,1] = stats.vals[13];
    err[i+1,1] = norm(x_ref[1:2] - x[1:2],Inf);

    (t,x,retcode,stats) = dop853(threebody,t0, T, x0, opt);
    if retcode != 1
        printFlag = false;
        break;
    end
    f_e[i+1,2] = stats.vals[13];
    err[i+1,2] = norm(x_ref[1:2] - x[1:2],Inf);

    (t,x,retcode,stats) = odex(threebody,t0, T, x0, opt);
    if retcode != 1
        printFlag = false;
        break;
    end
    f_e[i+1,3] = stats.vals[13];
    err[i+1,3] = norm(x_ref[1:2] - x[1:2],Inf);
end

# Save the plot in PNG format
if printFlag
    savePlotPNG("ArenstorfConvTest",f_e,err,solverNames);
else
    println("Cannot generate plot due to solver failure.")
end
end
end

```

3.3 RopePrecisionTest

```
# Load all the required packages
using ODEInterface
@ODEInterface.import_huge
loadODESolvers();

# Number of subdivisions of the rope
global n = 40;

# Define the system of ODEs
function rope(t,x,dx)
    n2 = n*n; # n^2
    n3by4 = convert{Int64,3*n/4}; # 3*n/4

    # Force in x-direction
    Fx = 0.4;
    # Force in y-direction
    Fy = cosh(4*t-2.5)^(-4);

    # Compute required matrices
    c = -cos(x[1:n-1]-x[2:n]);
    cDiag = [1;2*ones(n-2);3];
    C = spdiagsm((c,cDiag,c),(-1,0,1));

    d = -sin(x[1:n-1]-x[2:n]);
    D = spdiagsm((-d,d),(-1,1));

    # Compute the inhomogeneous term
    v = -(n2+n/2-n*[1:n;]).*sin(x[1:n])-n2*sin(x[1:n])*Fx;
    v[1:n3by4] = v[1:n3by4] + n2*cos(x[1:n3by4])*Fy;

    w = D*v+x[n+1:2*n].^2;
    u = C\w;

    # Write down the system
    dx[1:n] = x[n+1:2*n];
    dx[n+1:2*n] = C*v + D*u;

    return nothing
end

# Initial Conditions
t0 = 0.0; T = 3.723; x0=zeros(2*n);

# Compute the "reference solution"
opt = OptionsODE(OPT_EPS => 1.11e-16,OPT_RHS_CALLMODE => RHS_CALL_INSITU,
OPT_RTOL => 1e-16,OPT_ATOL=>1e-16);
(t,x_ref,retcode,stats) = dop853(rope,t0, T, x0, opt);

if retcode != 1
    println("Reference solution failed")
else
    # Initialization for the loop
    # f_e = function evaluations
    f_e = zeros{Int32,89,3};
```

```

# err = error for last step using infinity norm
err = zeros(Float64,89,3);

# solverNames = names of the solvers used for the plot
solverNames = ["DOPRI5","DOP853","ODEX"];

# Compute all the solutions
for i=0:88

    # Set up the tolerance
    Tol = 10^(-3-i/8);

    # Set up solver options
    opt = OptionsODE(OPT_EPS => 1.11e-16,OPT_RHS_CALLMODE => RHS_CALL_INSITU,
    OPT_RTOL => Tol,OPT_ATOL => Tol);

    # Solve using DOPRI5
    (t,x,retcode,stats) = dopri5(rope,t0, T, x0, opt);
    # Check if solver was successful
    if retcode != 1
        printFlag = false;
        break;
    end
    f_e[i+1,1] = stats.vals[13];
    err[i+1,1] = norm(x_accurate[1:n] - x[1:n],Inf);

    # Solve using DOP853
    (t,x,retcode,stats) = dop853(rope,t0, T, x0, opt);
    # Check if solver was successful
    if retcode != 1
        printFlag = false;
        break;
    end
    f_e[i+1,2] = stats.vals[13];
    err[i+1,2] = norm(x_accurate[1:n] - x[1:n],Inf);

    # Solve using ODEX
    (t,x,retcode,stats) = odex(rope,t0, T, x0, opt);
    # Check if solver was successful
    if retcode != 1
        printFlag = false;
        break;
    end
    f_e[i+1,3] = stats.vals[13];
    err[i+1,3] = norm(x_accurate[1:n] - x[1:n],Inf);
end

# Save the plot in PNG format
if printFlag
    savePlotPNG("RopeConvTest",f_e,err,solverNames);
else
    println("Cannot generate plot due to solver failure")
end
end

```


3.4 vdpolPrecisionTest

```
# Load all the required packages
using ODEInterface
using ForwardDiff
@ODEInterface.import_huge
loadODESolvers();

# Define the right-hand function for automatic differentiation
function vdpolAD(x)
    return [x[2], ((1-x[1]^2)*x[2]-x[1])*1e6]
end

# Define the system for the solver
function vdpol(t,x,dx)
    dx[:] = vdpolAD(x);
    return nothing
end

# Define the Jacobian function using AD
function getJacobian(t,x,J)
    J[:,:] = ForwardDiff.jacobian(vdpolAD,x);
    return nothing
end

# Flag to check whether plot is to be generated and saved or not
# Also checks if all solvers are successful
printFlag = true;

# Initial conditions
t0 = 0.0; T = [1.0:11.0;]; x0 = [2.0,0.0];

# Get "reference solution"
Tol = 1e-14;
# for Tol < 1e-14 we get the error "TOLERANCES ARE TOO SMALL"
opt = OptionsODE(OPT_EPS=>1.11e-16,OPT_RTOL=>Tol, OPT_ATOL=>Tol,
OPT_RHS_CALLMODE => RHS_CALL_INSITU,
OPT_JACOBI MATRIX => getJacobian);

# Store only the desired component
# Here, only the first component is desired
# The second component is the first derivative of the first component
# due to the fact that it is a second order system.
# Hence error will be taken over the first component only.
x_ref = Array{Float64}(11);

for i=1:11
    (t,x,retcode,stats) = seulex(vdpol,t0, T[i], x0, opt);
    if retcode!=1
        printFlag = false;
        break;
    end
    x_ref[i] = x[1];
end

if printFlag
```

```

# Store the solver names for plotting
solverNames = ["RADAU", "RADAU5", "SEULEX"];

# Initialize the variables for plots
# f_e = number of function evaluations
f_e = zeros(33,3);
# err = error wrt ref solution over all time steps and components
err = zeros(33,3);

for i =0:32

    # Set the tolerance for current run
    Tol = 10^(-2-i/4);

    # Set solver options
    opt = OptionsODE(OPT_EPS=>1.11e-16, OPT_ATOL=>Tol, OPT_RTOL=>Tol,
    OPT_RHS_CALLMODE => RHS_CALL_INSITU,
    OPT_JACOBI MATRIX=>getJacobian);

    # Restart the solution for each end time
    # to ensure a more accurate solution
    # compared to dense output

    # Solve using RADAU
    x_radau = Array{Float64}(11);
    for j=1:11
        (t,x,retcode,stats) = radau(vdpol,t0, T[j], x0, opt);
        # If solver fails do not continue further
        if retcode != 1
            println("Solver RADAU failed");
            printFlag = false;
            break;
        end
        x_radau[j] = x[1];
        f_e[i+1,1] = stats.vals[13];
    end
    # If solver fails do not continue further
    if !printFlag
        break;
    end
    err[i+1,1] = norm(x_radau-x_ref,Inf);

    # Solve using RADAU5
    x_radau5 = Array{Float64}(11);
    for j=1:11
        (t,x,retcode,stats) = radau5(vdpol,t0, T[j], x0, opt);
        # If solver fails do not continue further
        if retcode != 1
            println("Solver RADAU5 failed");
            printFlag = false;
            break;
        end
        x_radau5[j] = x[1];
        f_e[i+1,2] = stats.vals[13];
    end
end

```

```

end
# If solver fails do not continue further
if !printFlag
    break;
end
err[i+1,2] = norm(x_radau5-x_ref,Inf);

# Solve using SEULEX
x_seulex = Array{Float64}(11);
for j=1:11
    (t,x,retcode,stats) = seulex(vdpol,t0, T[j], x0, opt);
    # If solver fails do not continue further
    if retcode != 1
        println("Solver seulex failed");
        printFlag = false;
        break;
    end
    x_seulex[j] = x[1];
    f_e[i+1,3] = stats.vals[13];
end
# If solver fails do not continue further
if !printFlag
    break;
end
# Get the error over all the components and
err[i+1,3] = norm(x_seulex-x_ref,Inf);
end

# Save the plot in PNG format
# if all the solvers were successful
if printFlag
    savePlotPNG("vdpolPrecisionTest",f_e,err,solverNames);
else
    println("Plot cannot be generated due to failure");
end
else
    println("Plot cannot be generated due to failure of reference solution");
end
end

```

3.5 RoberPrecisionTest

```
# Load all the required packages
using ODEInterface
using ForwardDiff
@ODEInterface.import_huge
loadODESolvers();

# Define the right-hand function for Automatic Differentiation
function roberAD(x)
    return [-0.04*x[1]+1e4*x[2]*x[3],
            0.04*x[1]-1e4*x[2]*x[3]-3e7*(x[2])^2,
            3*10^7*(x[2])^2]
end

# Define the system for the solver
function rober(t,x,dx)
    dx[1] = -0.04*x[1]+1e4*x[2]*x[3];
    dx[2] = 0.04*x[1]-1e4*x[2]*x[3]-3e7*(x[2])^2;
    dx[3] = 3*10^7*(x[2])^2;
    return nothing
end

# Automatic Differentiation for a more general problem
function getJacobian(t,x,J)
    J[:,:] = ForwardDiff.jacobian(roberAD,x);
    return nothing
end

# Flag to check whether plot is to be generated and saved or not
# Also checks if all solvers are successful
printFlag = true;

# Initial conditions
t0 = 0.0; T = 10.^[0.0:11.0;]; x0=[1.0,0.0,0.0];

# Get "reference solution"
# TolMin < 1e-14 gives error "TOLERANCES ARE TOO SMALL"
Tol = 1e-14;
opt = OptionsODE(OPT_RHS_CALLMODE => RHS_CALL_INSITU,
    OPT_RTOL => Tol, OPT_ATOL=>Tol*1e-6,OPT_EPS => 1.11e-16,
    OPT_JACOBI MATRIX => getJacobian);

x_ref = Array{Float64}(12,3)

# Due to a bug in seulex solver for dense output,
# we restart the solution for each required output time
for i=1:12
    (t,x,retcode,stats) = seulex(rober,t0, T[i], x0, opt);
    # If solver fails do not continue further
    if retcode!=1
        printFlag=false;
        break;
    end
    x_ref[i,1] = x[1];
    x_ref[i,2] = x[2];
end
```

```

    x_ref[i,3] = x[3];
end

if printFlag
    # Store the solver names for plotting
    solverNames = ["RADAU", "RADAU5", "SEULEX"];

    # Initialize the variables for plots
    # err = error wrt ref solution over all time steps and components
    err = zeros(33,3);
    # f_e = number of function evaluations
    f_e = zeros(33,3);

    # Loop over all the tolerances
    for m = 0:32

        # Set the tolerance for current run
        Tol = 10^(-2-m/4);

        # Set solver options
        opt = OptionsODE(OPT_EPS=>1.11e-16, OPT_ATOL=>Tol*1e-6, OPT_RTOL=>Tol,
            OPT_RHS_CALLMODE => RHS_CALL_INSITU,
            OPT_JACOBI MATRIX=>getJacobian);

        # Restart the solution for each end time
        # to ensure a more accurate solution
        # compared to dense output

        # Solve using RADAU
        x_radau = Array{Float64}(12,3);
        for j=1:12
            (t,x,retcode,stats) = radau(rober,t0, T[j], x0, opt);
            # If solver fails do not continue further
            if retcode != 1
                println("Solver RADAU failed");
                printFlag = false;
                break;
            end
            x_radau[j,1] = x[1];
            x_radau[j,2] = x[2];
            x_radau[j,3] = x[3];
            f_e[m+1,1] = stats.vals[13];
        end
        # If solver fails do not continue further
        if !printFlag
            break;
        end
        err[m+1,1] = norm([norm(x_radau[:,1]-x_ref[:,1],Inf),
            norm(x_radau[:,2]-x_ref[:,2],Inf),
            norm(x_radau[:,3]-x_ref[:,3],Inf)],Inf);

        # Solve using RADAU5
        x_radau5 = Array{Float64}(12,3);
        for j=1:12

```

```

        (t,x,retcode,stats) = radau5(rober,t0, T[j], x0, opt);
        # If solver fails do not continue further
        if retcode != 1
            println("Solver RADAU5 failed");
            printFlag = false;
            break;
        end
        x_radau5[j,1] = x[1];
        x_radau5[j,2] = x[2];
        x_radau5[j,3] = x[3];
        f_e[m+1,2] = stats.vals[13];
    end
    # If solver fails do not continue further
    if !printFlag
        break;
    end
    err[m+1,2] = norm([norm(x_radau5[:,1]-x_ref[:,1],Inf),
        norm(x_radau5[:,2]-x_ref[:,2],Inf),
        norm(x_radau5[:,3]-x_ref[:,3],Inf)],Inf);

    # Solve using SEULEX
    x_seulex = Array{Float64}(12,3);
    for j=1:12
        (t,x,retcode,stats) = seulex(rober,t0, T[j], x0, opt);
        # If solver fails do not continue further
        if retcode != 1
            println("Solver seulex failed");
            printFlag = false;
            break;
        end
        x_seulex[j,1] = x[1];
        x_seulex[j,2] = x[2];
        x_seulex[j,3] = x[3];
        f_e[m+1,3] = stats.vals[13];
    end
    # If solver fails do not continue further
    if !printFlag
        break;
    end
    err[m+1,3] = norm([norm(x_seulex[:,1]-x_ref[:,1],Inf),
        norm(x_seulex[:,2]-x_ref[:,2],Inf),
        norm(x_seulex[:,3]-x_ref[:,3],Inf)],Inf);
end

if printFlag
    savePlotPNG("RoberPrecisionTest",f_e,err,solverNames);
else
    println("Plot cannot be generated due to failure");
end
else
    println("Plot cannot be generated due to failure of reference solution.");
end
end

```

References

- [1] Ernst Hairer, Syvert P. Nørsett, Gerhard Wanner. *Solving Ordinary Differential Equations I - Nonstiff Problems*. Springer-Verlag, Heidelberg, 1993.
- [2] Ernst Hairer, Gerhard Wanner. *Solving Ordinary Differential Equations II - Stiff and Differential-Algebraic Problems*. Springer-Verlag, Heidelberg, 1996.
- [3] <http://dcjones.github.io/Gadfly.jl/>
- [4] <https://github.com/luchr/ODEInterface.jl/>
- [5] <https://github.com/JuliaLang/ODE.jl>
- [6] <https://github.com/sonVishal/ODEInterface.jl/>
- [7] <http://www.unige.ch/~hairer/testset/testset.html>